



IoT·인공지능·빅데이터 개론 및 실습

빅데이터 배치, 대화형 질의,
스트림 분석

서울대학교 컴퓨터공학부
전병곤


Contents

- 1 빅데이터 스토리지
- 2 데이터 처리 일반
- 3 대화형 질의
- 4 스트림 처리

1 빅데이터 스토리지



10,000 Feet View of Big Data Systems



1 빅데이터 스토리지



(1) Data Storage

- Write and read data in a distributed storage
- Distributed file system
 - Google file system (GFS)
 - Hadoop distributed file system (HDFS):
an open source implementation of GFS

1 빅데이터 스토리지



(2) Hadoop Distributed File System (HDFS)

➤ File system interface: file, directory operations

➤ Optimized for big data processing


- Streaming read of large data
- Large, sequential writes that append data to files
- Large blocks (e.g., 128MB per block)
- Fault tolerance

➤ Master-worker architecture

1 빅데이터 스토리지

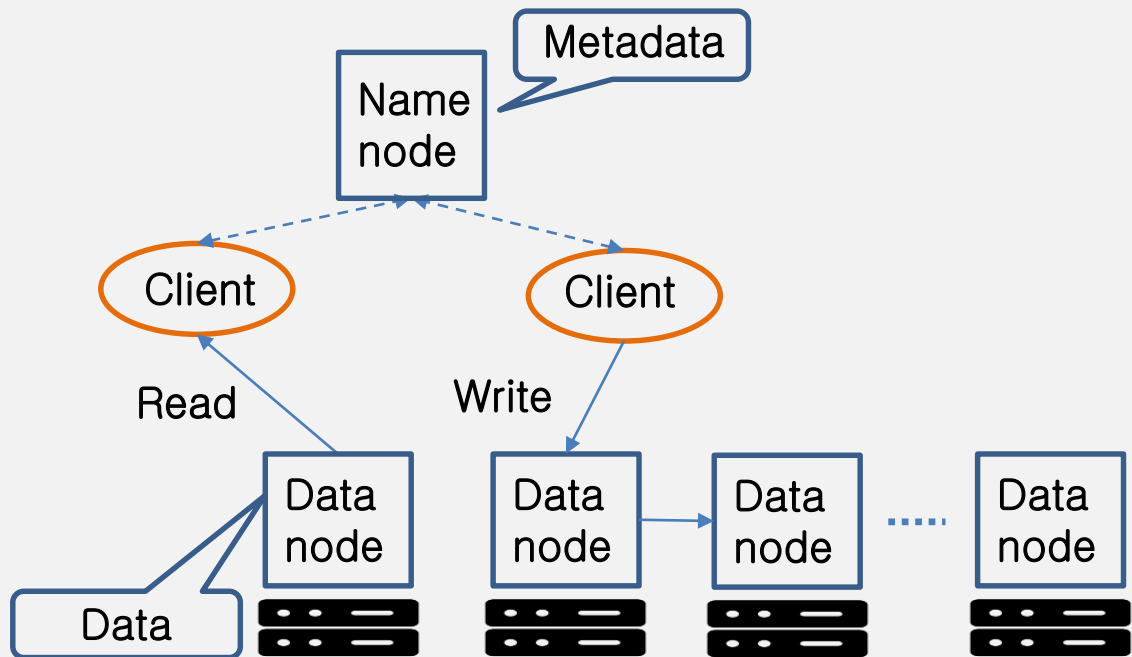
(2) Hadoop Distributed File System (HDFS)

➤ Master-worker architecture



1 빅데이터 스토리지

(2) Hadoop Distributed File System (HDFS)





(2) Hadoop Distributed File System (HDFS)

➤ HDFS architecture


- Data flow is decoupled from control flow
- Clients interact with the namenode for metadata operations
- Clients interact directly with datanodes for all files operations
- Performance can be improved by scheduling expensive data flow based on the network topology

Contents

- 1 빅데이터 스토리지
- 2 데이터 처리 일반
- 3 대화형 질의
- 4 스트림 처리

2 데이터 처리 일반

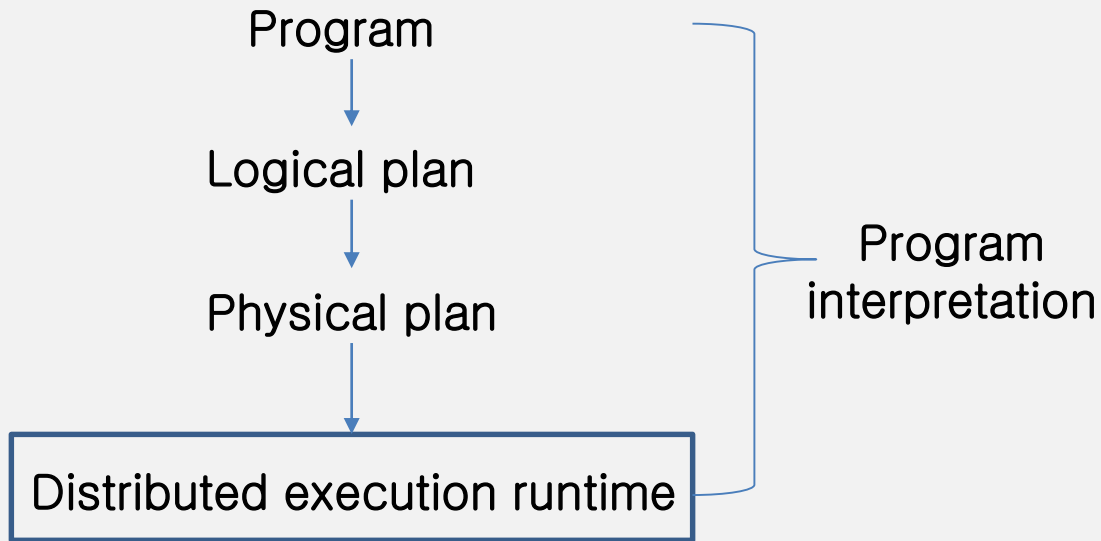
10,000 Feet View of Big Data Systems



2 데이터 처리 일반

(1) 데이터플로우 모델


➤ Data processing



2 데이터 처리 일반

(1) 데이터플로우 모델

➤ Data processing programming model




2 데이터 처리 일반

(1) 데이터플로우 모델

➤ Dataflow graph

- A Directed Acyclic Graph of operators
 - Vertex: operator
 - Edge: data dependency





(2) 맵리듀스

➤ MapReduce programming model

- Map function: process a key/value pair to generate a set of intermediate key/value pairs
- Reduce function: merges all intermediate values associated with the same intermediate key

(2) 맵리듀스


➤ MapReduce WordCount

- Wordcount problem: compute how many times each word appears in a collection of documents
- Map: $\text{lambda word: (word, 1)}$
- Reduce: $\text{lambda a, b: a + b}$

2 데이터 처리 일반

(2) 맵리듀스

➤ MapReduce : logical plan

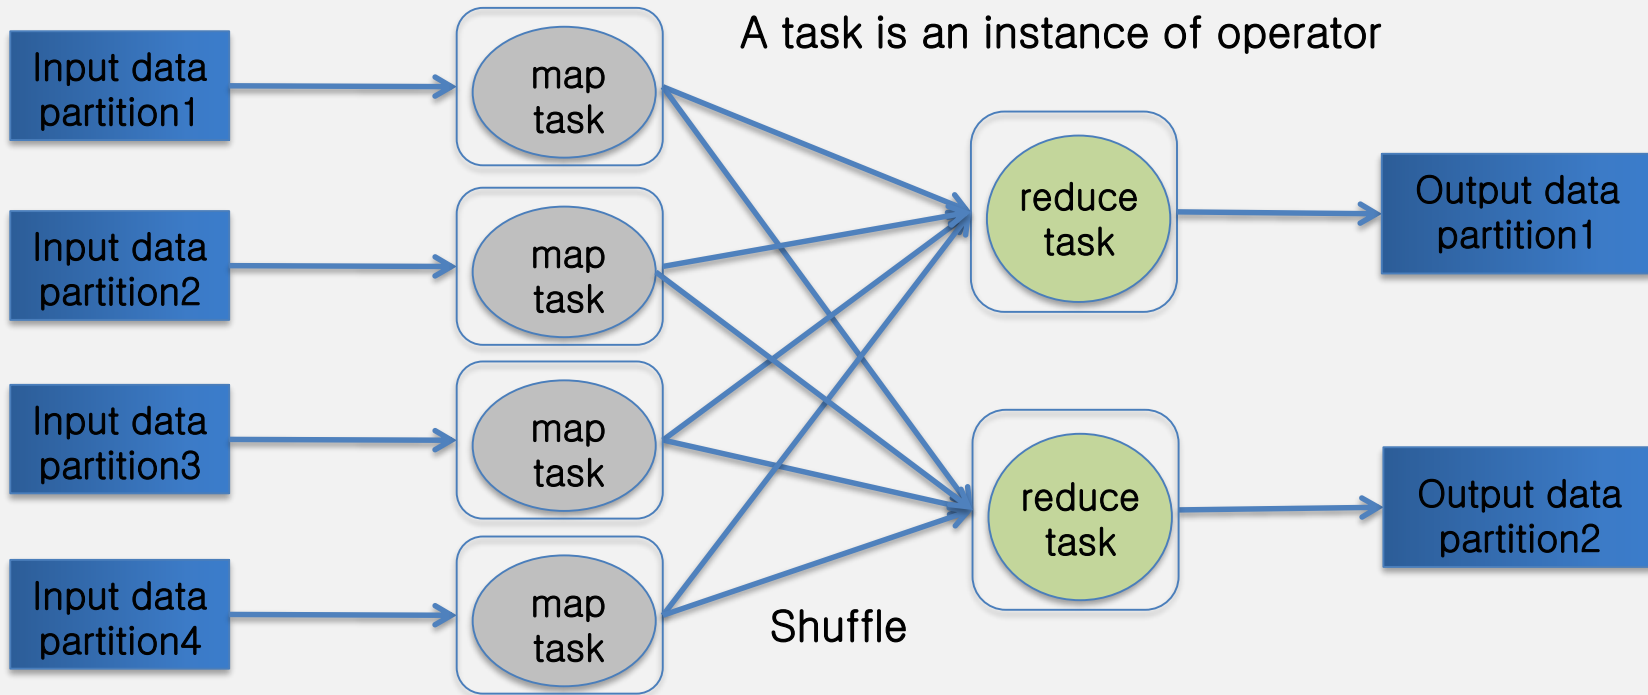


lambda word: (word, 1) lambda a, b: a + b

2 데이터 처리 일반

(2) 맵리듀스

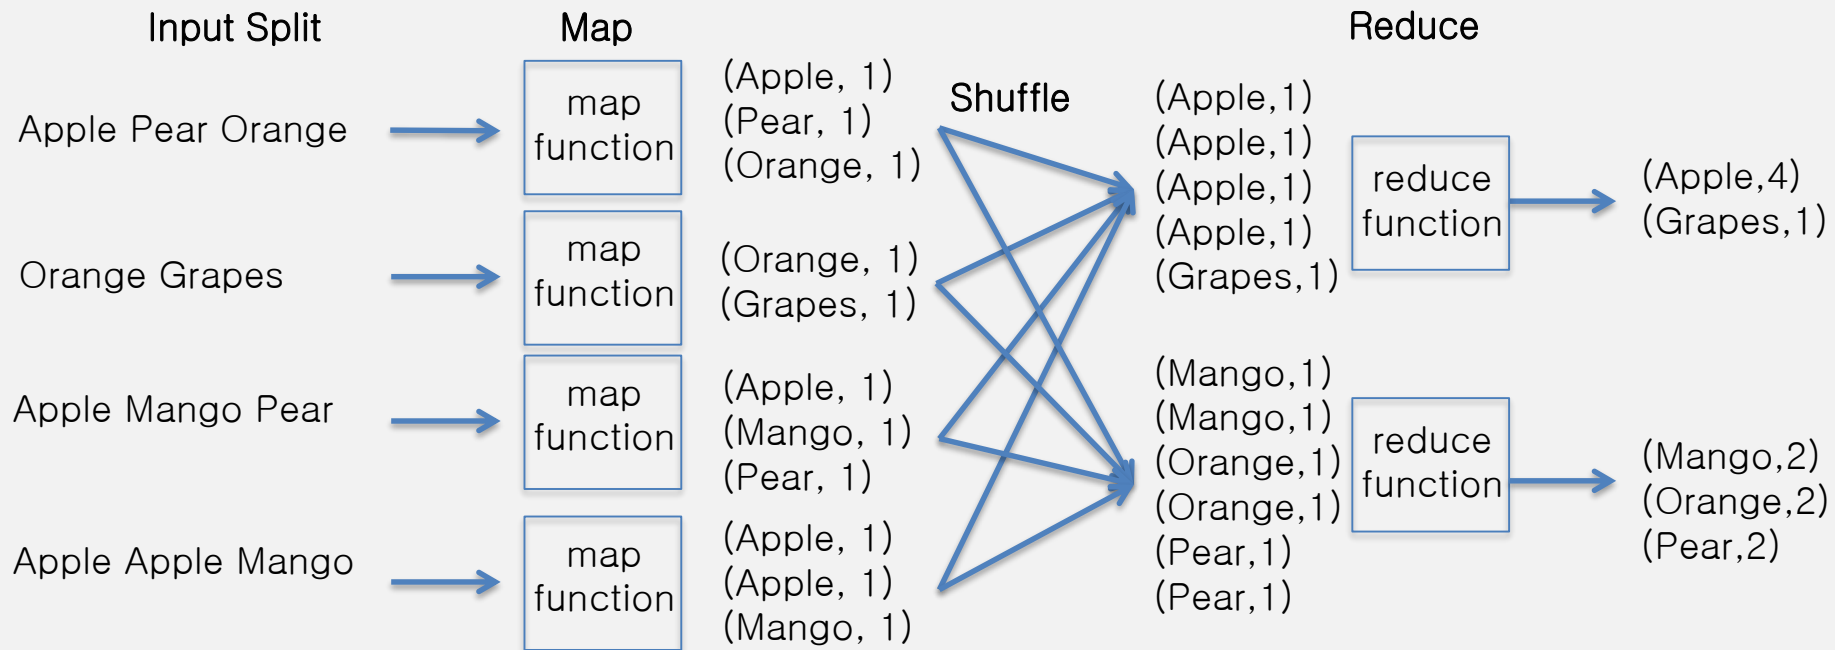
➤ MapReduce : physical plan



2 데이터 처리 일반

(2) 맵리듀스

➤ MapReduce WordCount execution example





(2) 맵리듀스

➤ MapReduce runtime handles

- Scheduling
- Fault Tolerance
- Scalability
- Elasticity



(2) 맵리듀스

➤ Limitations of MapReduce

- Limited programming flexibility
 - In particular, inefficient for multi-pass algorithms
- No efficient primitives for data sharing
 - State between MapReduce passes goes through distributed file systems

2 데이터 처리 일반

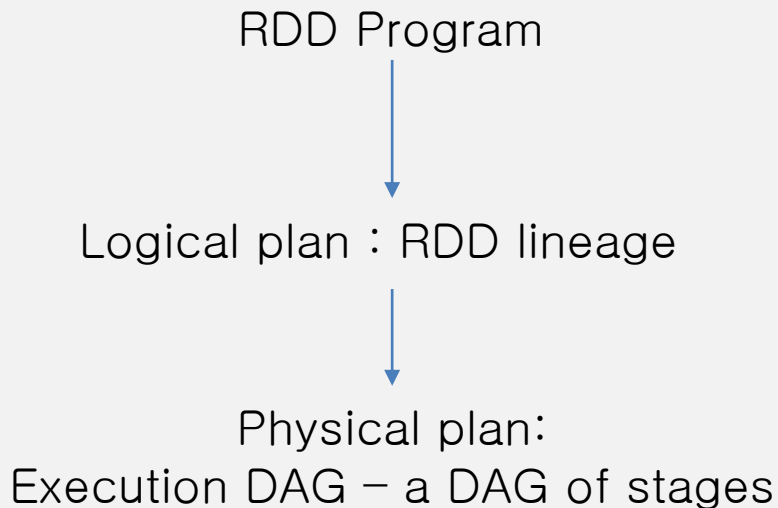
(3) 스파크

- Generic dataflow graph: a DAG of operators
- In-memory computing with Resilient Distributed Datasets (RDDs)
- Easy-to-use programming interface: transformation on RDDs

2 데이터 처리 일반

(3) 스파크

➤ Spark logical plan, physical plan



(3) 스파크

➤ RDDs

- Collections of “immutable” objects spread across a cluster
- Statically typed: `RDD[T]` has objects of type `T`
- Built through parallel transformations
- All transformations are lazy
- Automatically rebuilt on a failure through lineage
- Two types of operations
 - Transformations (e.g., `map`, `filter`, `groupBy`)
 - Actions (e.g., `count`, `collect`)
- RDD can be persisted into storage in memory or disk

2 데이터 처리 일반

(3) 스파크

➤ RDD lineage example

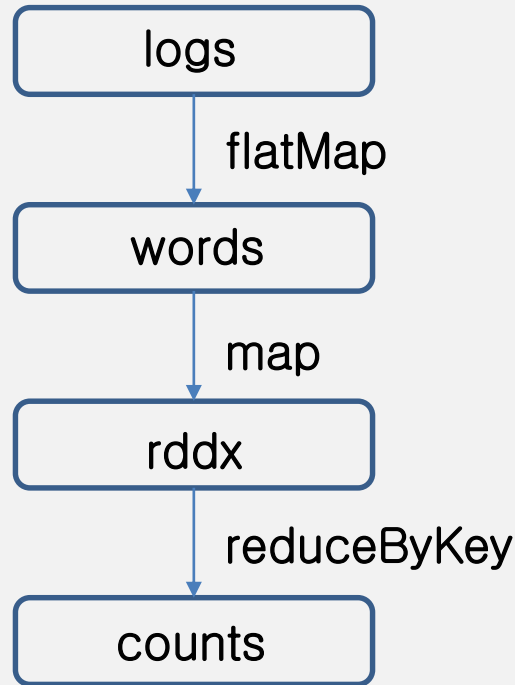
```
logs = sc.textFile("hdfs://data/logs")
```

```
words = logs.flatMap(lambda line: line.split())
```

```
counts = words.map(lambda word: (word, 1))
```

```
                .reduceByKey(lambda c1, c2: c1 +
```


```
c2)  
counts.saveAsTextFile("hdfs://result/counts")
```



2 데이터 처리 일반

(3) 스파크

➤ Execution DAG





(3) 스파크

➤ Batch processing

- Data preprocessing
- Extract-Transform-Load (ETL)
- Data processing to do in bulk
(not latency-critical)

Contents

- 1 빅데이터 스토리지
- 2 데이터 처리 일반
- 3 대화형 질의
- 4 스트림 처리

3 대화형 질의

(1) Interactive Query

- Structured data
- Use SQL to express queries on structured data
- SQL: a domain-specific language used for managing data in relational models

3 대화형 질의

(2) Structured Data

- Data managed in a table
- A table consists of rows and columns

Table name: customer

id	name	age	Place
1	Alex Kim	32	Seoul
2	Jane Lee	25	Seoul
3	Matt Park	23	Incheon

3 대화형 질의

(2) Structured Data

➤ Operations on structured data

- Projection: select specified columns from a table
 - E.g., $\Pi_{\{\text{name}, \text{age}\}}(\text{customer})$

name	age
Alex Kim	32
Jane Lee	25
Matt Park	23

3 대화형 질의

(2) Structured Data

➤ Operations on structured data

- Selection: select a subset of rows from a table that meets a predicate
 - E.g., $\sigma_{\text{age} > 30}(\text{customer})$

id	name	age	place
1	Alex Kim	32	Seoul

3 대화형 질의

(2) Structured Data

➤ Operations on structured data

- Aggregation: aggregate on a column
 - E.g., $G_{\text{sum}(\text{place})}$ (customer)

place	count
Seoul	2
Incheon	1

(2) Structured Data

➤ Operations on structured data

- Join: compute rows in two tables that are equal on their common column names
 - Natural join
 - Left outer join
 - Right outer join

3 대화형 질의

(2) Structured Data

➤ Operations on structured data

- Natural join
 - E.g., customer ⋈ department

Table name: customer

id	name	age	place
1	Alex Kim	32	Seoul
2	Jane Lee	25	Seoul
3	Matt Park	23	Incheon

Table name: department

place	lead
Seoul	Bob
Incheon	Alice

id	name	age	place	lead
1	Alex Kim	32	Seoul	Bob
2	Jane Lee	25	Seoul	Bob
3	Matt Park	23	Incheon	Alice

←
customer ⋈ department

(3) SparkSQL: Interactive Query in Spark

- DataFrame: a Dataset organized into named columns
- Conceptually equivalent to a table in a relational database or a data frame in R/Python but with richer optimizations under the hood
- Can be constructed from various sources such as structured data files, tables in Hive, external databases, or existing RDDs
- Support relational operations
- An SQL query is translated into RDD transformations to be executed

Contents

- 1 빅데이터 스토리지
- 2 데이터 처리 일반
- 3 대화형 질의
- 4 스트림 처리

4 스트림 처리

(1) Stream Processing

➤ Stream data

- Unbounded data
- Event time \neq data processing time

➤ Real-time processing

- Window operation

4 스트림 처리

(1) Stream Processing

➤ What: stream processing model

➤ How: stream processing execution engine

- Continuous query: e.g., Storm, Heron, Flink
- Micro-batch: e.g., SparkStreaming


(2) Streaming Patterns

- Element-wise transformations
- Processing-time based windows
- Event-time based windows
- Session windows


4 스트림 처리

(3) Windowing


Fixed



Sliding



Sessions

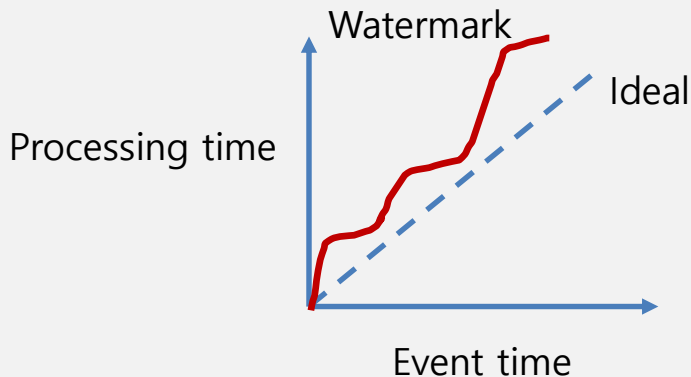


Time

4 스트림 처리

(4) Watermarks

- Watermark: no timestamp earlier than the watermark will be seen

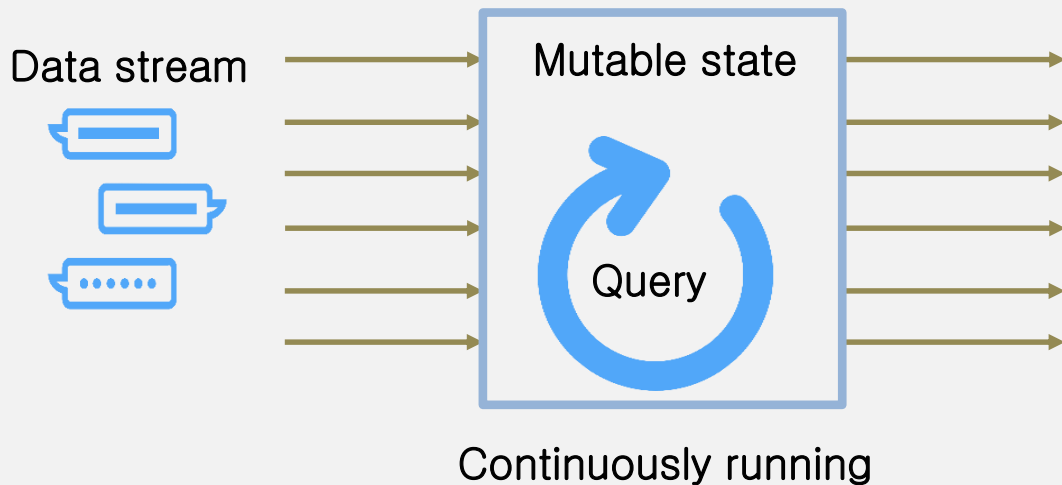


- Watermark too slow: results delayed
- Watermark too fast: some data late

4 스트림 처리

(5) Stream Query Execution Model


➤ Continuous Query: e.g., Storm, Heron, Flink



4 스트림 처리

(5) Stream Query Execution Model

- Micro-batch: e.g., SparkStreaming
- Dissociate computation from state: make state immutable and break computation into small, deterministic, stateless tasks



▼ 빅데이터 스토리지

▼ 데이터 처리 일반

▼ 대화형 질의

▼ 스트림 처리