



# IoT·인공지능·빅데이터 개론 및 실습

그래프 처리,  
머신러닝/딥러닝

서울대학교 컴퓨터공학부  
전병곤

# Contents

- 1 그래프 처리
- 2 머신러닝/딥러닝
- 3 머신러닝/딥러닝 분산학습

# 1 그래프 처리


## (1) 빅데이터 분석 방법

- 배치 분석
- 대화형 질의
- 스트림 처리
- 그래프 처리
- 머신 러닝/딥러닝

# 1 그래프 처리

## (2) Graph Data

Web graph



[출처]: <http://labs.criteo.com/2014/05/web-graph-seen-criteo/>

Social graph




[출처]: <http://blog.shuttlecloud.com/more-than-the-social-graph/>

# 1 그래프 처리

## (2) Graph Definition and Algorithms

### ➤ Graph G definition

- $V$ 
  - A set of vertices
  - $L(v)$  – value of vertex  $v$
- $E$ 
  - A set of edges (directed or undirected edges)
  - $W(u, v)$  – value of edge  $(u, v)$



### ➤ Variations of the above definition

# 1 그래프 처리

## (2) Graph Definition and Algorithms

### ➤ Graph algorithms

- Shorted paths
- Bipartite matching
- Connected components
- Random walk
- PageRank
- ...

# 1 그래프 처리

## (2) Graph Definition and Algorithms

### ➤ Graph algorithms

- Triangle counting
- Community detection
- Motif finding
- Social circles
- ...

## (3) Graph Processing Model

### ➤ Why not MapReduce or MPI for graph processing?

- MapReduce: not efficient
- MPI: implementing specific graph processing algorithms  
→ not reusable for other graph processing algorithms



## (3) Graph Processing Model

### ➤ Think like a vertex

- e.g., Google Pregel, Apache Giraph
- Vertex as a computation unit

### ➤ Think like a (sub)graph

- A subgraph as a computation unit

## (3) Graph Processing Model

### ➤ Think like a vertex

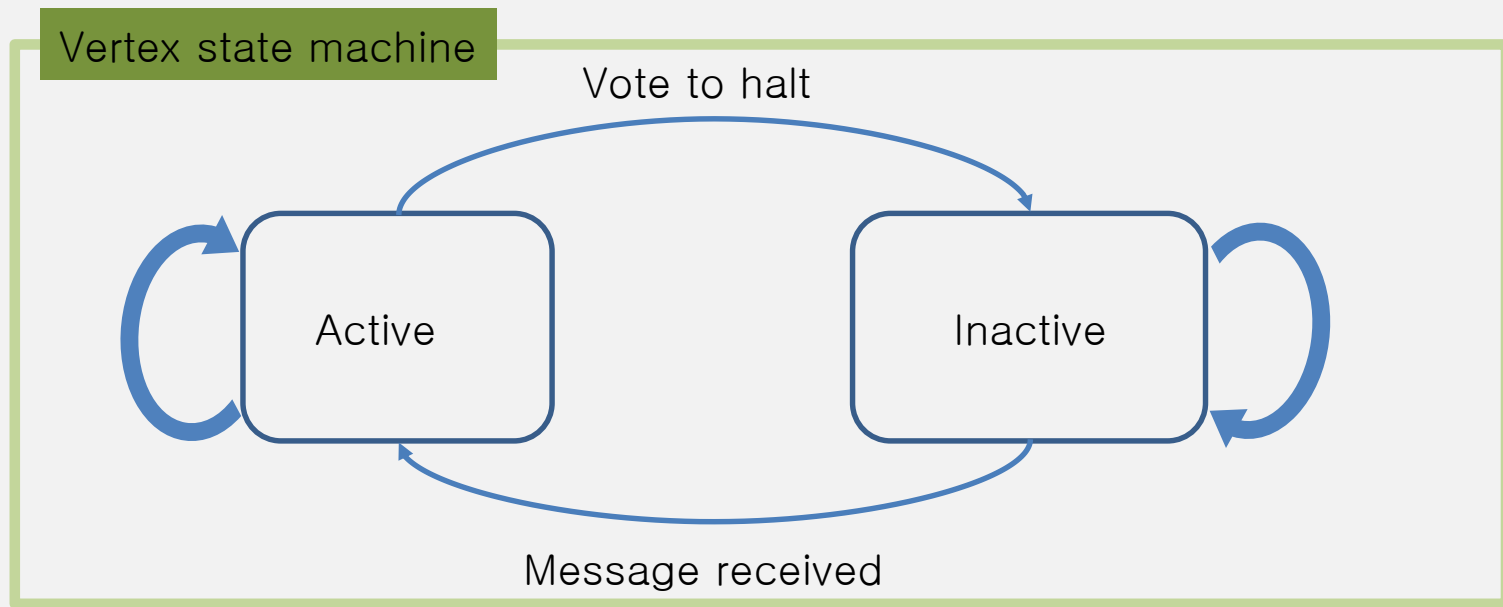
- Computations consist of a sequence of iterations called supersteps (i.e., bulk synchronous parallel)
- During a superstep, the framework invokes a user-defined function for each vertex in parallel
- A vertex can
  - read messages sent to  $V$  in superstep  $S-1$
  - modify the state of  $V$  and its outgoing edges
  - send messages to other vertices that will be received at superstep  $S+1$

# 1 그래프 처리

## (3) Graph Processing Model

➤ Think like a vertex (e.g., Google Pregel, Apache Giraph)

- Algorithm termination is based on every vertex voting to halt




# 1 그래프 처리

## (3) Graph Processing Model

### ➤ Example of the think-like-a-vertex model

- Compute the maximum vertex value of a graph




Superstep 0

# 1 그래프 처리

## (3) Graph Processing Model

### ➤ Example of the think-like-a-vertex model

- Compute the maximum vertex value of a graph




# 1 그래프 처리

## (3) Graph Processing Model

### ➤ Example of the think-like-a-vertex model

- Compute the maximum vertex value of a graph




# 1 그래프 처리

## (3) Graph Processing Model

### ➤ Example of the think-like-a-vertex model


- Compute the maximum vertex value of a graph



# 1 그래프 처리

## (4) Distributed Graph Processing

- Master-worker architecture
- Master divides a graph into partitions(subgraphs) to maximize parallelism





## (4) Distributed Graph Processing

- Master initiates a superstep
- Each Worker runs vertex computation for the vertexes in the partition
  - Sends and receives messages across machine boundaries
- Workers notify Master their state changes
- Master checks if the current state meets its termination condition
  - If so, terminates the processing
  - If not, initiates the next superstep

# Contents

- 1 그래프 처리
- 2 머신러닝/딥러닝
- 3 머신러닝/딥러닝 분산학습



# (1) Classical Machine Learning vs Deep Learning

## ➤ Classical machine learning

- Human-crafted features
- Great fit for data mining applications

## ➤ Deep learning

- No feature extraction
- Great fit for hard vision, speech, language problems



### (2) Applications

#### ➤ Classical machine learning(ML) applications

- Regression
- Classification
- Clustering
- Topic modeling
- Collaborative filtering
- Frequent pattern mining
- Ranking
- ...



## (2) Applications

### ➤ Deep learning(DL) applications

- Image classification
- Speech-to-text, text-to-speech
- Video understanding
- Image, video style transfer
- Text understanding – dialog
- Machine translation
- Ranking ads, feeds, news
- Robot control, game play (e.g., GO)
- Self-driving car
- ...

### (3) ML(DL) Workflow


#### ➤ Two main steps of ML

- Training
  - Constructs a machine learning model by optimizing loss given training data
- Inference
  - Uses a trained model to predict for new input data

## 2 머신러닝/딥러닝

### (3) ML(DL) Workflow

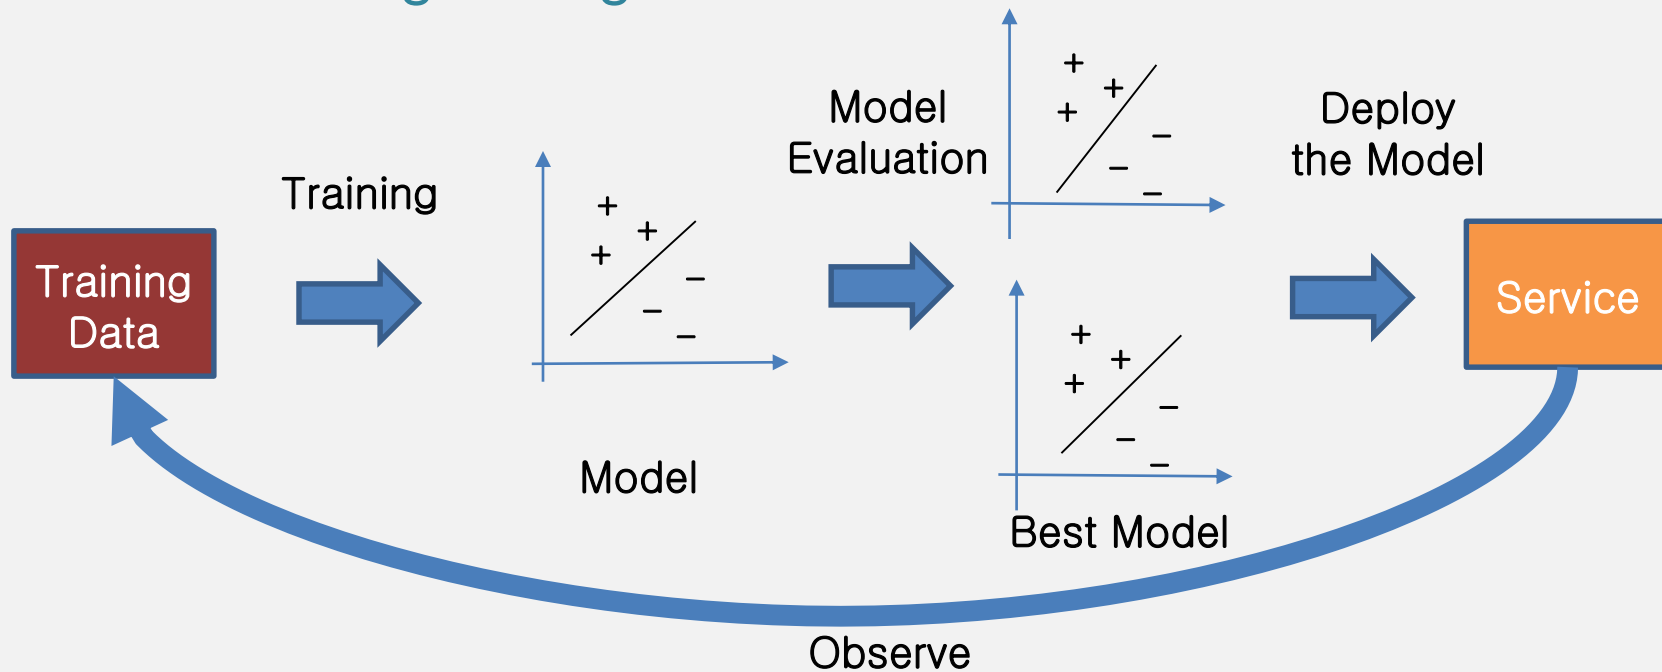
#### ➤ Machine learning workflow



## 2 머신러닝/딥러닝

### (3) ML(DL) Workflow

#### ➤ Machine learning training workflow





## 2 머신러닝/딥러닝

### (4) ML(DL) Framework

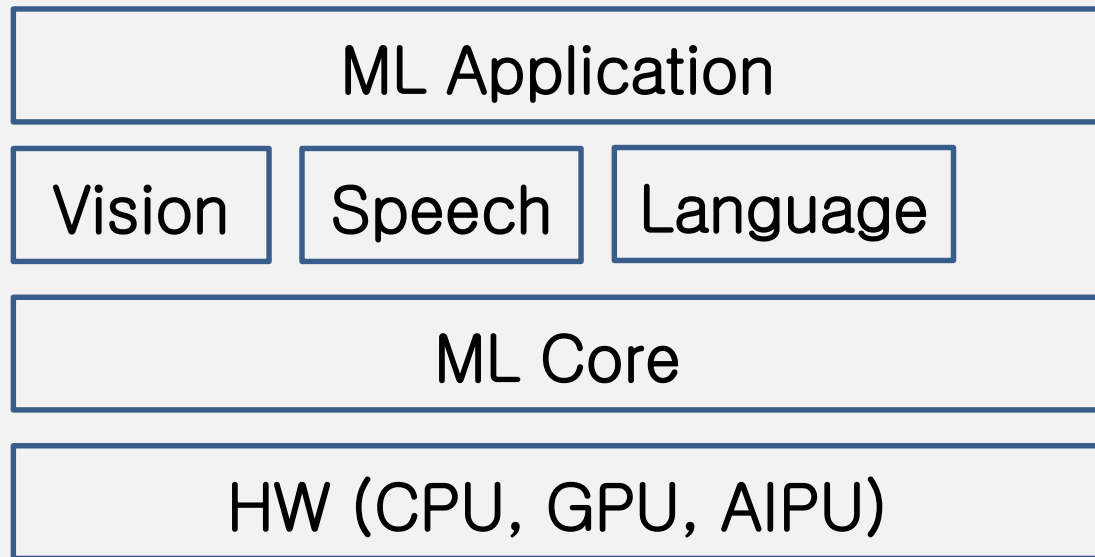
➤ Machine learning (deep learning) framework



## 2 머신러닝/딥러닝

### (4) ML(DL) Framework

#### ➤ Machine learning software stack



### (4) ML(DL) Framework

#### ➤ GPU (Graphics Processing Unit)

- Originally for graphics processing
- Around 2004, GPUs were programmable “enough” to do some non-graphics computations
  - Severely limited by graphics programming model (shader programming)
- In 2006, GPUs became “fully” programmable
  - NVIDIA releases “CUDA” language to write non-graphics programs that will run on GPUs
- Great fit for deep learning

### (4) ML(DL) Framework

#### ➤ ML core high-level structure

- Python frontend: define machine learning models (mostly neural nets)
- C++ backend: execute defined models
  - CPU
  - GPU(s)
  - AIPU (e.g., TPU)
  - Distributed multi-machine

### (4) ML(DL) Framework

#### ➤ ML core main ideas

- Describe mathematical computation
- Represent computation as a computation graph
- Execute the graph with the given data batch
- Training: iterate execution with massive dataset to optimize a goal
  - The framework supports auto-differentiation
- Inference: execute once with a data input to predict

### (4) ML(DL) Framework

#### ➤ Framework types: symbolic vs imperative

- Symbolic: define-and-run
  - e.g., TensorFlow, Caffe2, MXNet
  - Better to optimize
  - Easier to deploy
  - Harder to program and debug
- Imperative: define-by-run
  - e.g., PyTorch, TensorFlow eager, MXNet imperative
  - Easier to program and debug
  - Slower than symbolic
  - Harder to deploy

### (4) ML(DL) Framework

#### ➤ Google TensorFlow

- Symbolic ML framework
- Express numerical computation as a computation graph
  - Node: operation which has any number of inputs and outputs
  - Edge: tensor which flow between nodes
- Tensor: N-dimensional array
  - 1-dimension: Vector
  - 2-dimension: Matrix
  - E.g., image represented as 3-d tensor rows, cols, color
- Tensors flow through the graph => TensorFlow

### (4) ML(DL) Framework

#### ➤ Google TensorFlow

- The graph's compiled to CPU / GPU / AIPU code
- Salient features
  - Fine-grained ops
  - Dynamic control flow: condition, loop
  - Persistent state maintenance/update



### (4) ML(DL) Framework

#### ➤ TensorFlow programming model

- Operation: abstract computation  
(e.g., matrix multiply)
- Kernel: a particular implementation of an operation  
that can be run on a particular type of device  
(e.g., CPU or GPU)
- Variable: returns a handle to a persistent mutable  
tensor that survives across executions of a graph.
- Session: runs a graph

## 2 머신러닝/딥러닝

### (4) ML(DL) Framework

#### ➤ TensorFlow code example


```
import tensorflow as tf
```

```
b = tf.Variable(tf.zeros((100,)))
```

```
W = tf.Variable(tf.random_uniform((784, 100), -1, 1))
```

```
x = tf.placeholder(tf.float32, (None, 784))
```

```
h_i = tf.nn.relu(tf.matmul(x, W) + b)
```



For training, TensorFlow automatically creates its backpropagation graph

### (4) ML(DL) Framework

#### ➤ FaceBook PyTorch

- Imperative ML framework
- Express numerical computation like numpy
- Programming model
  - Operation : abstract computation (e.g., matrix multiply)
  - Tensor : imperative N-dimensional array but runs either on CPU or GPU
  - Variable : wrapper of Tensor. It constructs a chain of operations between the tensors, so that the gradients can flow back

### (4) ML(DL) Framework

#### ➤ Interoperability between frameworks : Open Neural Network Exchange (ONNX)

- An open source format for AI models
- An extensible computation graph model and definitions of built-in operators and standard data types
- Caffe2, PyTorch, Cognitive Toolkit, MXNet support ONNX

# Contents

- 1 그래프 처리
- 2 머신러닝/딥러닝
- 3 머신러닝/딥러닝 분산학습



### (1) Distributed training issues

- Parallelism: data parallelism, model parallelism, hybrid parallelism
- Synchronization: synchronous, asynchronous, bounded synchronous
- Synchronization architecture: parameter server architecture, MPI



### (1) Distributed training issues

#### ➤ Parallelism

- Data Parallelism
  - $D = D_1 \cup D_2 \cup \dots \cup D_n$
  - Each worker  $i$  processes  $M$  with  $D_i$ .
- Model Parallelism
  - $M = M_1 \cup M_2 \cup \dots \cup M_n$
  - Each worker  $i$  processes  $M_i$  with  $D$ .
- Hybrid Parallelism



## (1) Distributed training issues

### ➤ Synchrony of training

- Synchronous training
- Asynchronous training
- Bounded-synchronous training






### (1) Distributed training issues

#### ➤ Synchronization architecture

- Parameter server
- Message passing interface (MPI)

### ③ 머신러닝/딥러닝 분산학습


#### (2) Parameter Server



### (2) Parameter Server

#### ➤ Server


- Maintains a partition of the globally shared parameters
- Performs global aggregation steps



### (2) Parameter Server

#### ➤ Worker

- Performs computation with (a portion of) training data communicates with servers
- Updating and retrieving the shared parameters





### (2) Parameter Server

#### ➤ Parameter-server-based synchronous training

- Step 1 : Workers compute gradients with training data and push them to servers
- Step 2 : Each Server receives gradients from Workers, aggregates them, and applies the sums to update the model parameters
- Step 3 : Workers pull the new model parameters
- The above steps iterate until training converges

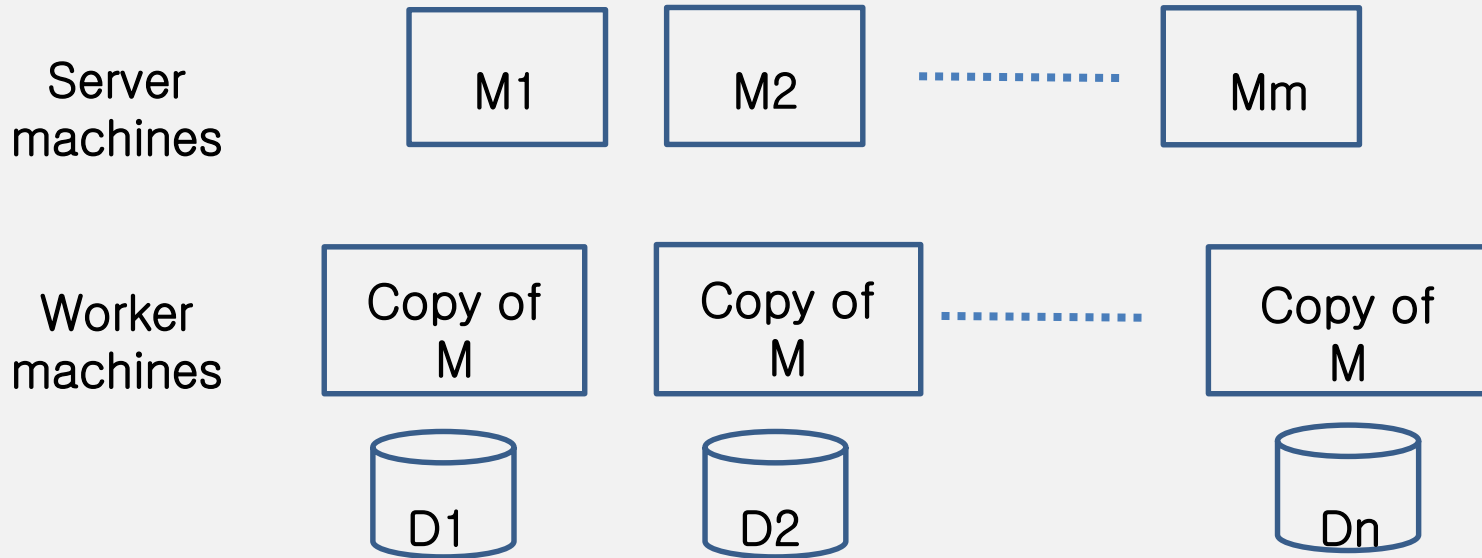
### 3 머신러닝/딥러닝 분산학습



## (2) Parameter Server

### ➤ Parallelism


- Data Parallelism




### 3 머신러닝/딥러닝 분산학습



#### (3) MPI Collective Communication



Broadcast




Scatter

### 3 머신러닝/딥러닝 분산학습




## (3) MPI Collective Communication



Gather

AllGather – Gather + Broadcast



Reduce

AllReduce – Reduce + Broadcast





### (3) MPI Collective Communication

#### ➤ MPI-based synchronous training

- Step 1 : Workers compute gradients with training data
- Step 2 : Workers run AllReduce (or AllGather) to aggregate them and apply the sums to update the model parameters
- The above steps iterate until training converges

- ▼ 그래프 처리
- ▼ 머신러닝/딥러닝
- ▼ 머신러닝/딥러닝 분산학습